

Vijay(P7) 3/08/2023

In [327]:	1 import numpy as np 2 import pandas as pd 3 import seaborn as sns 4 import matplotlib.pyplot as plt
In [328]:	1 df=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\madrid_2007.csv") 2 df
Out[328]:	
	225120 rows × 17 columns

In [329]:	1 df=df.dropna()
In [330]:	1 df.columns 2
Out[330]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')	

In [331]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25443 entries, 4 to 225119
Data columns (total 17 columns):
 #   Column   Non-Null Count  Dtype  
---  --  
 0   date      25443 non-null   object 
 1   BEN        25443 non-null   float64
 2   CO         25443 non-null   float64
 3   EBE        25443 non-null   float64
 4   MXY        25443 non-null   float64
 5   NMHC       25443 non-null   float64
 6   NO_2       25443 non-null   float64
 7   NOx        25443 non-null   float64
 8   OXY        25443 non-null   float64
 9   O_3        25443 non-null   float64
 10  PM10       25443 non-null   float64
 11  PM25       25443 non-null   float64
 12  PXY        25443 non-null   float64
 13  SO_2       25443 non-null   float64
 14  TCH        25443 non-null   float64
 15  TOL        25443 non-null   float64
 16  station    25443 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```

In [332]: 1 data=df[['BEN', 'CO','station']]
2 data
3

Out[332]:

	BEN	CO	station
4	4.64	1.86	28079006
21	1.98	0.31	28079024
25	2.82	1.42	28079099
30	4.65	1.89	28079006
47	1.97	0.30	28079024
...
225073	2.12	0.47	28079006
225094	0.87	0.45	28079099
225098	0.95	0.41	28079006
225115	0.30	0.45	28079024
225119	0.53	0.40	28079099

25443 rows × 3 columns

In [333]:

```
1 df=df.head(10000)
2 df
```

Out[333]:

		date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2
4		2007-12-01 01:00:00	4.64	1.86	4.26	7.98	0.57	145.100006	573.900024	3.49	52.689999	106.500000	15.900000	3.56	40.230000
21		2007-12-01 01:00:00	1.98	0.31	2.56	6.06	0.35	76.059998	208.899994	1.70	1.000000	37.799999	25.580000	1.78	11.310000
25		2007-12-01 01:00:00	2.82	1.42	3.15	7.02	0.49	123.099998	402.399994	2.60	7.160000	70.809998	37.009998	2.67	25.670000
30		2007-12-01 02:00:00	4.65	1.89	4.41	8.21	0.65	151.000000	622.700012	3.55	58.080002	117.099998	17.049999	3.57	36.459999
47		2007-12-01 02:00:00	1.97	0.30	2.15	5.08	0.33	78.760002	189.800003	1.62	1.000000	34.740002	24.730000	1.59	10.500000
...	
89387		2007-11-22 20:00:00	0.49	0.23	0.69	1.69	0.21	47.919998	52.080002	0.47	20.950001	8.260000	5.890000	0.41	7.670000
89391		2007-11-22 20:00:00	1.22	0.61	1.63	3.12	0.26	83.550003	149.300003	1.25	8.220000	32.700001	19.360001	1.23	14.150000
89396		2007-11-22 21:00:00	3.47	1.03	2.51	4.89	0.31	74.779999	192.000000	2.25	4.390000	34.529999	12.030000	2.18	18.150000
89413		2007-11-22 21:00:00	0.68	0.25	0.66	1.64	0.24	75.330002	95.309998	0.46	4.140000	15.440000	11.040000	0.40	9.050000
89417		2007-11-22 21:00:00	1.34	0.60	1.86	3.27	0.24	80.660004	141.699997	1.36	8.110000	18.049999	12.630000	1.29	15.120000

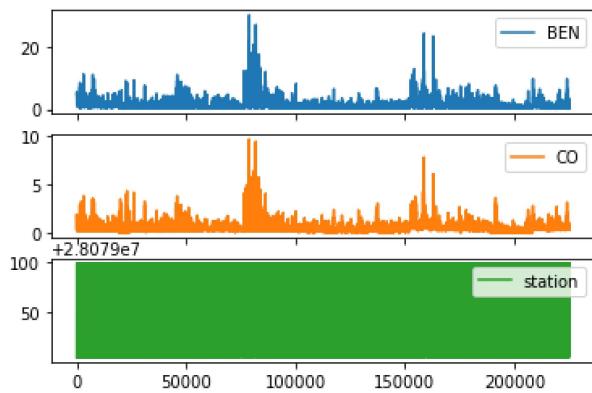
10000 rows × 17 columns

In [334]:

```
1 data.plot.line(subplots=True)
```

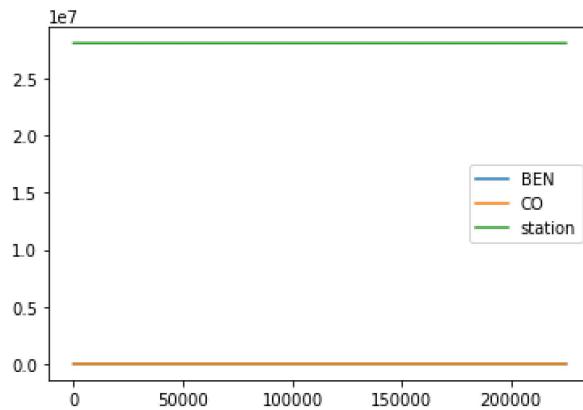
Out[334]:

```
array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



```
In [335]: 1 data.plot.line()  
2
```

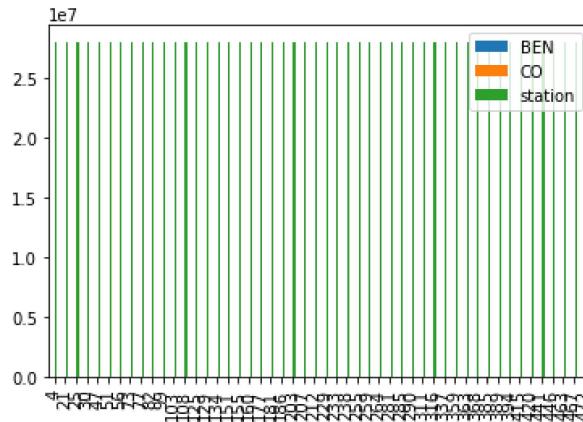
Out[335]: <AxesSubplot:>



```
In [336]: 1 b=data[0:50]  
2
```

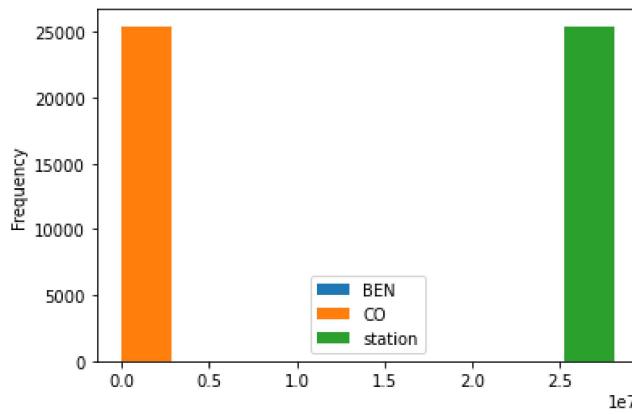
```
In [337]: 1 b.plot.bar()
```

Out[337]: <AxesSubplot:>



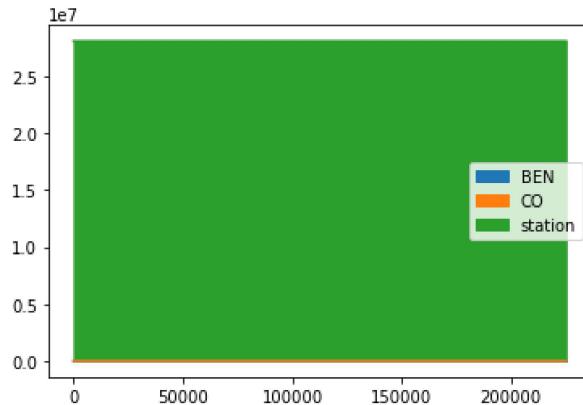
```
In [338]: 1 data.plot.hist()  
2
```

Out[338]: <AxesSubplot:ylabel='Frequency'>



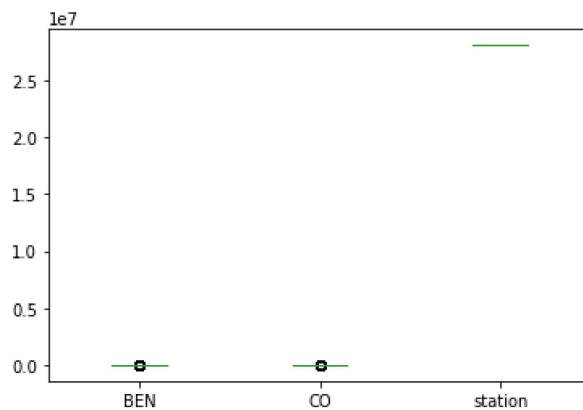
```
In [339]: 1 data.plot.area()
```

```
Out[339]: <AxesSubplot:>
```



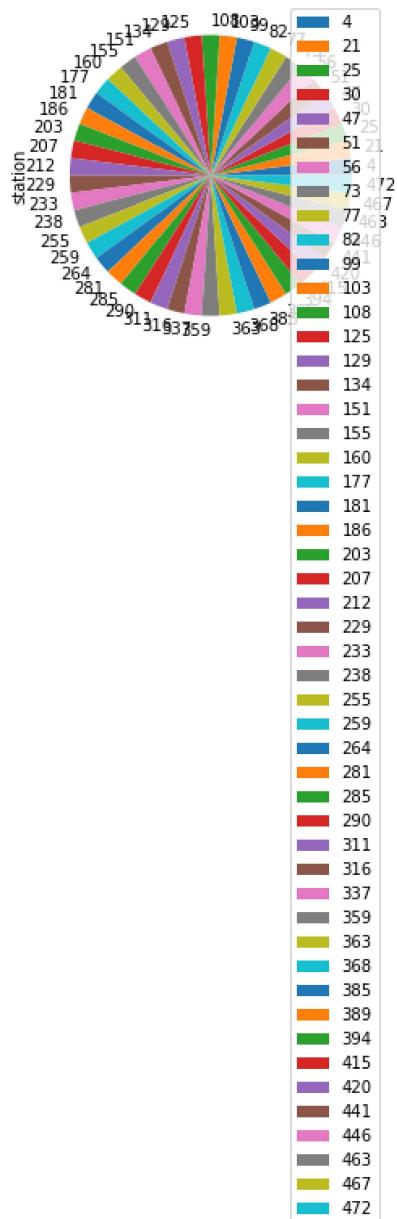
```
In [340]: 1 data.plot.box()
```

```
Out[340]: <AxesSubplot:>
```



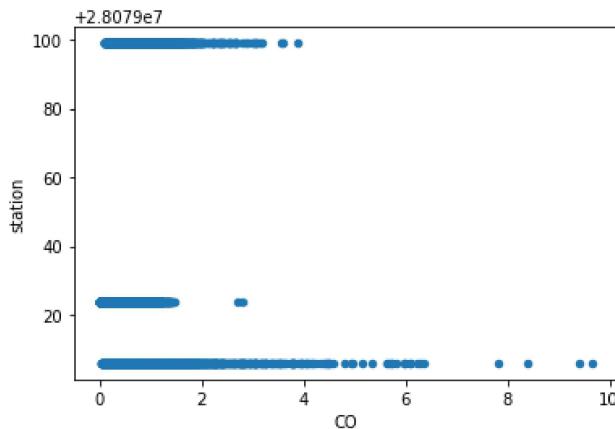
In [341]: 1 b.plot.pie(y='station')

Out[341]: <AxesSubplot:ylabel='station'>



```
In [342]: 1 data.plot.scatter(x='CO' ,y='station')
2
```

Out[342]: <AxesSubplot:xlabel='CO', ylabel='station'>



```
In [343]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10000 entries, 4 to 89417
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        10000 non-null   object 
 1   BEN          10000 non-null   float64
 2   CO           10000 non-null   float64
 3   EBE          10000 non-null   float64
 4   MXY          10000 non-null   float64
 5   NMHC         10000 non-null   float64
 6   NO_2          10000 non-null   float64
 7   NOx          10000 non-null   float64
 8   OXY          10000 non-null   float64
 9   O_3           10000 non-null   float64
 10  PM10         10000 non-null   float64
 11  PM25         10000 non-null   float64
 12  PXY          10000 non-null   float64
 13  SO_2          10000 non-null   float64
 14  TCH           10000 non-null   float64
 15  TOL           10000 non-null   float64
 16  station       10000 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 1.4+ MB
```

```
In [344]: 1 df.describe()
2
```

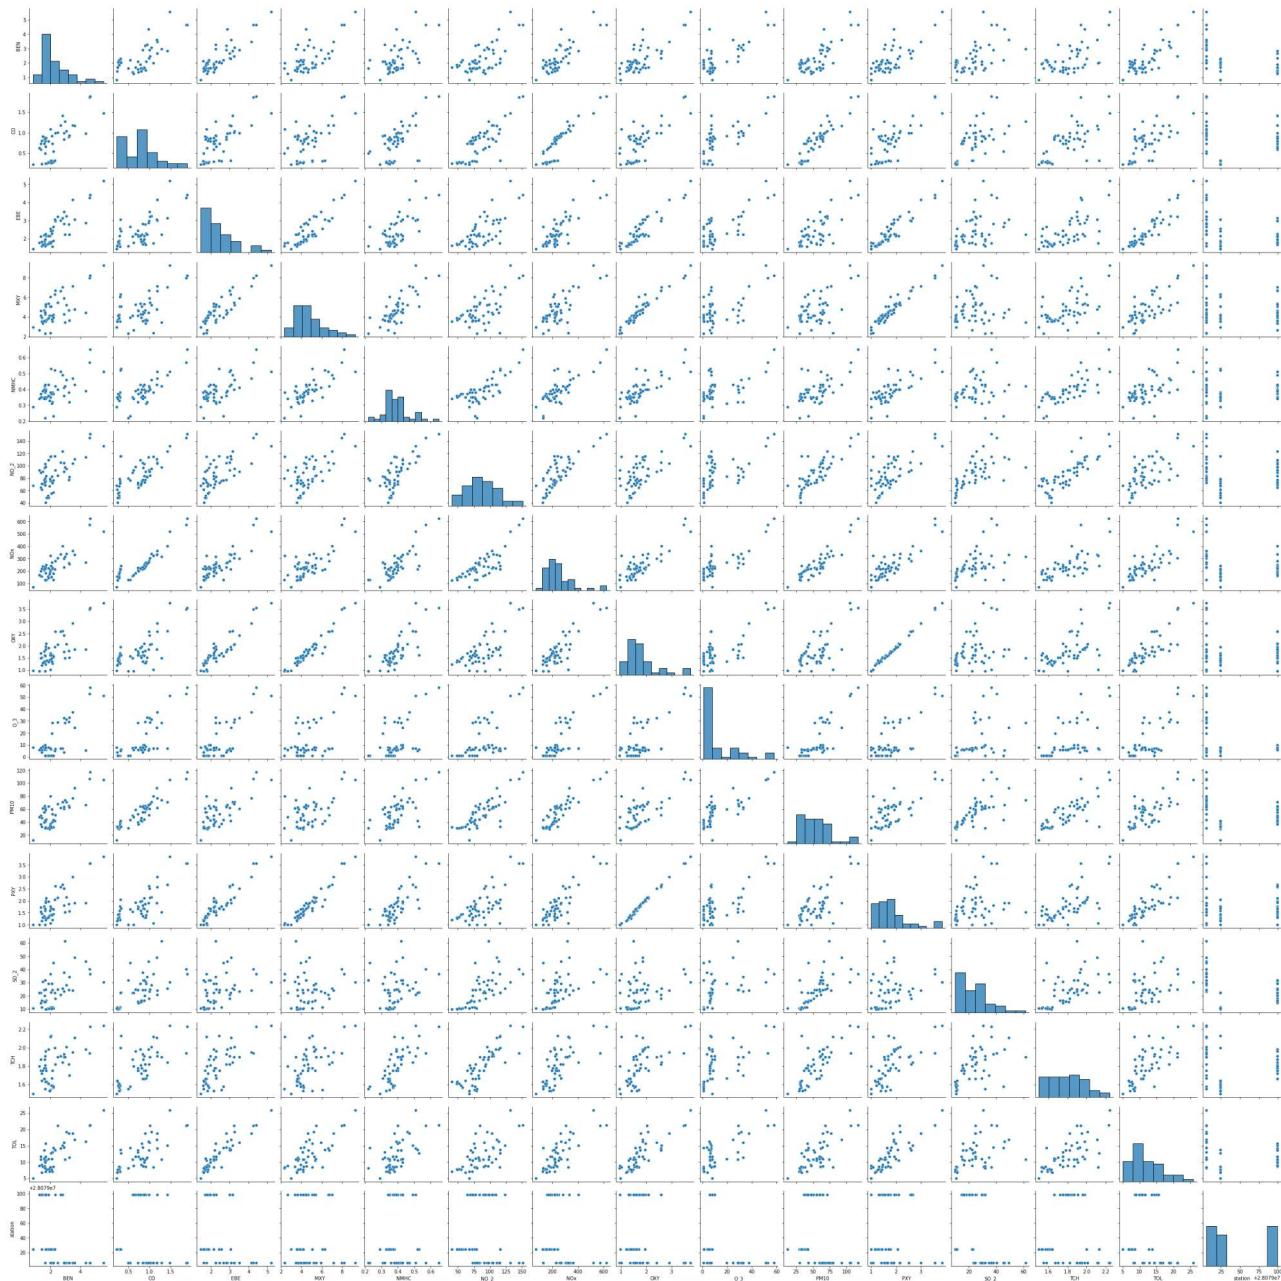
Out[344]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	SO_2
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	1.337737	0.520786	1.508041	2.654044	0.315455	65.391101	129.601638	1.303739	29.1
std	1.465736	0.491611	1.292714	2.677514	0.162820	38.943317	129.665140	1.129495	24.1
min	0.150000	0.010000	0.190000	0.170000	0.000000	2.510000	2.620000	0.110000	0.1
25%	0.560000	0.250000	0.790000	1.000000	0.210000	37.919998	49.892501	0.730000	7.9
50%	0.930000	0.375000	1.140000	1.850000	0.280000	60.525000	95.309998	1.000000	22.4
75%	1.620000	0.630000	1.830000	3.362500	0.400000	86.162502	168.325001	1.510000	44.4
max	30.139999	9.660000	23.379999	46.730000	2.570000	430.299988	1847.000000	32.540001	140.

```
In [345]: 1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2   'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [346]: 1 sns.pairplot(df1[0:50])
```

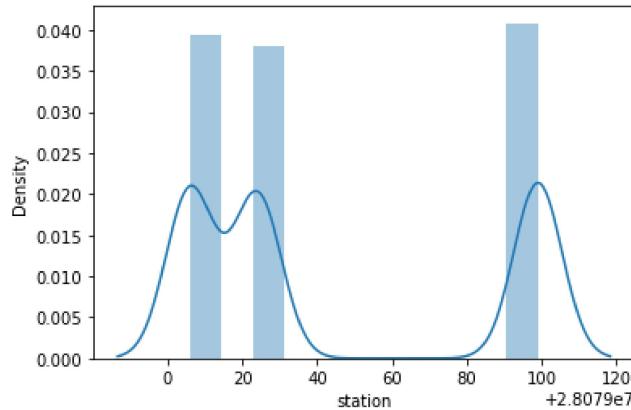
```
Out[346]: <seaborn.axisgrid.PairGrid at 0x21019f38280>
```



```
In [347]: 1 sns.distplot(df1['station'])
2
```

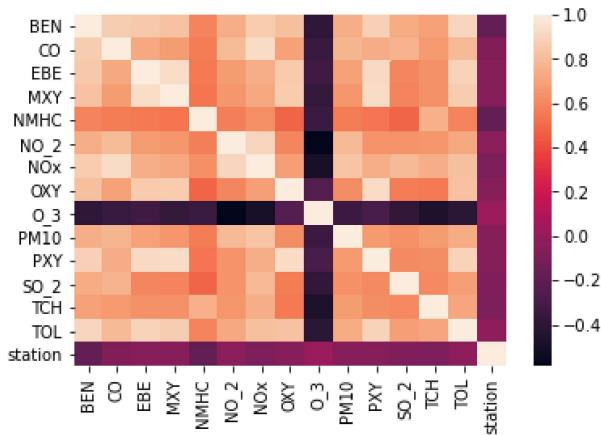
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
`warnings.warn(msg, FutureWarning)

```
Out[347]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [348]: 1 sns.heatmap(df1.corr())
```

```
Out[348]: <AxesSubplot:>
```



```
In [349]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
4
```

```
In [350]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
3
```

```
In [351]: 1 from sklearn.linear_model import LinearRegression
2 lr=LinearRegression()
3 lr.fit(x_train,y_train)
4
```

```
Out[351]: LinearRegression()
```

```
In [352]: 1 lr.intercept_
```

```
Out[352]: 28079024.11153279
```

```
In [353]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
2 coeff
```

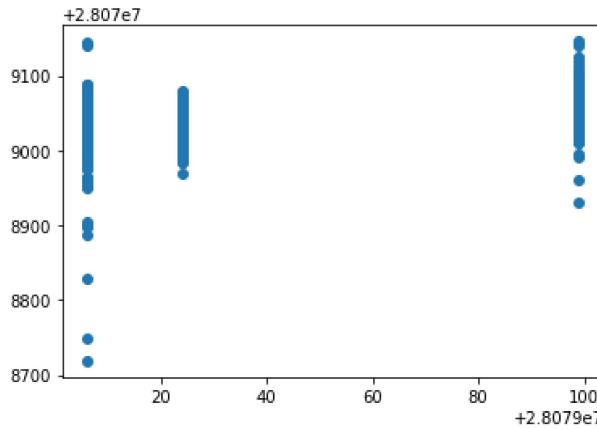
Out[353]:

Co-efficient

BEN	-39.715434
CO	38.232571
EBE	1.649810
MXY	-1.727348
NMHC	-67.799240
NO_2	0.119086
NOx	-0.048293
OXY	0.569914
O_3	-0.094167
PM10	0.010241
PXY	14.261016
SO_2	0.261422
TCH	21.922055
TOL	3.920577

```
In [354]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

Out[354]: <matplotlib.collections.PathCollection at 0x2102b1257f0>



```
In [355]: 1 lr.score(x_test,y_test)
2
```

Out[355]: 0.18832072724652216

```
In [356]: 1 lr.score(x_train,y_train)
2
```

Out[356]: 0.24039322600259538

```
In [357]: 1 from sklearn.linear_model import Ridge,Lasso
2
```

```
In [358]: 1 rr=Ridge(alpha=10)
2 rr.fit(x_train,y_train)
3
```

Out[358]: Ridge(alpha=10)

```
In [359]: 1 rr.score(x_test,y_test)
2
```

Out[359]: 0.18956540856303727

```
In [360]: 1 rr.score(x_train,y_train)
2
```

Out[360]: 0.23975318959885694

```
In [361]: 1 la=Lasso(alpha=10)
2 la.fit(x_train,y_train)
```

Out[361]: Lasso(alpha=10)

```
In [362]: 1 la.score(x_train,y_train)
2
```

Out[362]: 0.013157253655736056

```
In [363]: 1 la.score(x_test,y_test)
2
```

Out[363]: 0.008479502920505033

```
In [364]: 1 from sklearn.linear_model import ElasticNet
2 en=ElasticNet()
3 en.fit(x_train,y_train)
4
```

Out[364]: ElasticNet()

```
In [365]: 1 en.coef_
2
```

Out[365]: array([-9.58857946e+00, 0.00000000e+00, -0.00000000e+00, -1.36724997e-01,
-8.21327427e-01, 1.48081172e-01, -5.05636798e-02, 0.00000000e+00,
2.06137262e-02, -2.32329912e-03, 0.00000000e+00, 2.98477369e-02,
-0.00000000e+00, 1.90310519e+00])

```
In [366]: 1 en.intercept_
2
```

Out[366]: 28079041.88418573

```
In [367]: 1 prediction=en.predict(x_test)
2
```

```
In [368]: 1 en.score(x_test,y_test)
```

Out[368]: 0.08889326711786572

```
In [369]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
5
```

36.24640511702622
1512.82603329351
38.89506438217464

```
In [370]: 1 from sklearn.linear_model import LogisticRegression
2
```

```
In [371]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df['station']
4
```

```
In [372]: 1 feature_matrix.shape
2
```

Out[372]: (10000, 14)

```
In [373]: 1 target_vector.shape
2
```

Out[373]: (10000,)

```
In [374]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [375]: 1 fs=StandardScaler().fit_transform(feature_matrix)
2
```

```
In [376]: 1 logr=LogisticRegression(max_iter=10000)
2 logr.fit(fs,target_vector)
```

Out[376]: LogisticRegression(max_iter=10000)

```
In [377]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
2
```

```
In [378]: 1 prediction=logr.predict(observation)
2 print(prediction)
3
```

[28079099]

```
In [379]: 1 logr.classes_
2
```

Out[379]: array([28079006, 28079024, 28079099], dtype=int64)

```
In [380]: 1 logr.score(fs,target_vector)
2
```

Out[380]: 0.8101

```
In [381]: 1 logr.predict_proba(observation)[0][0]
2
```

Out[381]: 1.2204737088913592e-28

```
In [382]: 1 logr.predict_proba(observation)
2
```

Out[382]: array([[1.22047371e-28, 4.19009794e-17, 1.00000000e+00]])

```
In [383]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [384]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

Out[384]: RandomForestClassifier()

```
In [385]: 1 parameters={'max_depth':[1,2,3,4,5],  
2   'min_samples_leaf':[5,10,15,20,25],  
3   'n_estimators':[10,20,30,40,50]  
4 }
```

```
In [386]: 1 from sklearn.model_selection import GridSearchCV  
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")  
3 grid_search.fit(x_train,y_train)  
4
```

```
Out[386]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param_grid={'max_depth': [1, 2, 3, 4, 5],  
'min_samples_leaf': [5, 10, 15, 20, 25],  
'n_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')
```

```
In [387]: 1 grid_search.best_score_
```

```
Out[387]: 0.8162857142857143
```

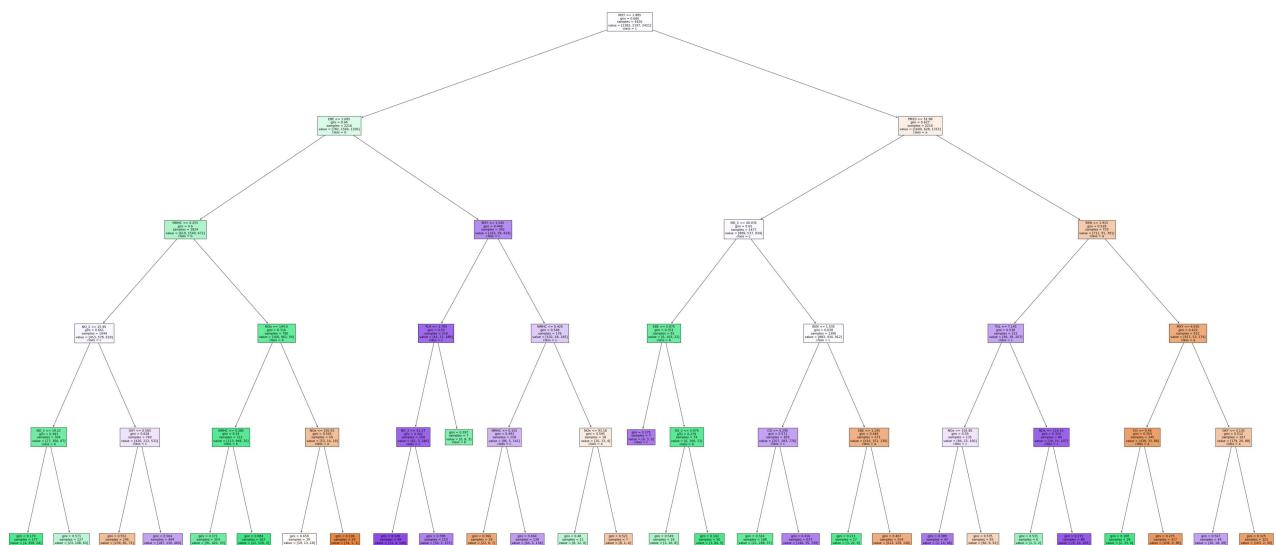
```
In [388]: 1 rfc_best=grid_search.best_estimator_
```

In [389]:

```
1 from sklearn.tree import plot_tree
2 plt.figure(figsize=(80,40))
3 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

```
Out[389]: [Text(2192.142857142857, 1993.2, 'MXY <= 1.885\ngini = 0.666\nsamples = 4426\nvalue = [2382, 2197, 242
1]\nnclass = c'),
Text(1175.7857142857142, 1630.8000000000002, 'EBE <= 1.005\ngini = 0.64\nsamples = 2216\nvalue = [782,
1569, 1106]\nnclass = b'),
Text(637.7142857142857, 1268.4, 'NMHC <= 0.255\ngini = 0.6\nsamples = 1824\nvalue = [619, 1540, 672]\n
class = b'),
Text(318.85714285714283, 906.0, 'NO_2 <= 25.95\ngini = 0.661\nsamples = 1044\nvalue = [453, 578, 618]
\nnclass = c'),
Text(159.42857142857142, 543.5999999999999, 'NO_2 <= 19.22\ngini = 0.383\nsamples = 304\nvalue = [27,
366, 87]\nnclass = b'),
Text(79.71428571428571, 181.1999999999982, 'gini = 0.179\nsamples = 177\nvalue = [4, 258, 24]\nnclass
= b'),
Text(239.1428571428571, 181.1999999999982, 'gini = 0.571\nsamples = 127\nvalue = [23, 108, 63]\nnclass
= b'),
Text(478.2857142857142, 543.5999999999999, 'OXY <= 0.565\ngini = 0.628\nsamples = 740\nvalue = [426, 2
12, 531]\nnclass = c'),
Text(398.57142857142856, 181.1999999999982, 'gini = 0.552\nsamples = 246\nvalue = [239, 82, 71]\nclas
s = a'),
Text(558.0, 181.1999999999982, 'gini = 0.564\nsamples = 494\nvalue = [187, 130, 460]\nnclass = c'),
Text(956.5714285714284, 906.0, 'NOx <= 144.0\ngini = 0.316\nsamples = 780\nvalue = [166, 962, 54]\nncla
ss = b'),
Text(797.1428571428571, 543.5999999999999, 'NMHC <= 0.385\ngini = 0.24\nsamples = 721\nvalue = [113, 9
48, 35]\nnclass = b'),
Text(717.4285714285713, 181.1999999999982, 'gini = 0.371\nsamples = 354\nvalue = [91, 422, 33]\nnclass
= b'),
Text(876.8571428571428, 181.1999999999982, 'gini = 0.084\nsamples = 367\nvalue = [22, 526, 2]\nnclass
= b'),
Text(1116.0, 543.5999999999999, 'NOx <= 192.55\ngini = 0.545\nsamples = 59\nvalue = [53, 14, 19]\nclas
s = a'),
Text(1036.2857142857142, 181.1999999999982, 'gini = 0.658\nsamples = 39\nvalue = [19, 13, 18]\nnclass
= a'),
Text(1195.7142857142856, 181.1999999999982, 'gini = 0.106\nsamples = 20\nvalue = [34, 1, 1]\nnclass = a'),
Text(1713.8571428571427, 1268.4, 'MXY <= 1.545\ngini = 0.449\nsamples = 392\nvalue = [163, 29, 434]\nnc
lass = c'),
Text(1514.5714285714284, 906.0, 'TCH <= 1.705\ngini = 0.33\nsamples = 216\nvalue = [61, 11, 289]\nncla
ss = c'),
Text(1434.8571428571427, 543.5999999999999, 'NO_2 <= 51.17\ngini = 0.302\nsamples = 209\nvalue = [61,
3, 286]\nnclass = c'),
Text(1355.142857142857, 181.1999999999982, 'gini = 0.149\nsamples = 99\nvalue = [11, 2, 149]\nnclass = c
'),
Text(1514.5714285714284, 181.1999999999982, 'gini = 0.398\nsamples = 110\nvalue = [50, 1, 137]\nnclass
= c'),
Text(1594.2857142857142, 543.5999999999999, 'gini = 0.397\nsamples = 7\nvalue = [0, 8, 3]\nnclass =
b'),
Text(1913.1428571428569, 906.0, 'NMHC <= 0.425\ngini = 0.548\nsamples = 176\nvalue = [102, 18, 145]\nnc
lass = c'),
Text(1753.7142857142856, 543.5999999999999, 'NMHC <= 0.155\ngini = 0.493\nsamples = 158\nvalue = [86,
5, 141]\nnclass = c'),
Text(1673.999999999998, 181.1999999999982, 'gini = 0.366\nsamples = 19\nvalue = [22, 0, 7]\nnclass = a
'),
Text(1833.4285714285713, 181.1999999999982, 'gini = 0.464\nsamples = 139\nvalue = [64, 5, 134]\nnclass
= c'),
Text(2072.5714285714284, 543.5999999999999, 'NOx <= 93.18\ngini = 0.595\nsamples = 18\nvalue = [16,
13, 4]\nnclass = a'),
Text(1992.8571428571427, 181.1999999999982, 'gini = 0.48\nsamples = 11\nvalue = [8, 12, 0]\nnclass =
b'),
Text(2152.285714285714, 181.1999999999982, 'gini = 0.521\nsamples = 7\nvalue = [8, 1, 4]\nnclass =
a'),
Text(3208.499999999995, 1630.8000000000002, 'PM10 <= 51.98\ngini = 0.627\nsamples = 2210\nvalue = [16
00, 628, 1315]\nnclass = a'),
Text(2590.7142857142853, 1268.4, 'NO_2 <= 40.035\ngini = 0.65\nsamples = 1477\nvalue = [889, 537, 934]
\nnclass = c'),
Text(2311.7142857142853, 906.0, 'EBE <= 0.875\ngini = 0.351\nsamples = 81\nvalue = [6, 103, 22]\nnclass
= b'),
Text(2232.0, 543.5999999999999, 'gini = 0.375\nsamples = 7\nvalue = [0, 3, 9]\nnclass = c'),
Text(2391.428571428571, 543.5999999999999, 'SO_2 <= 5.975\ngini = 0.279\nsamples = 74\nvalue = [6, 10
0, 13]\nnclass = b'),
Text(2311.7142857142853, 181.1999999999982, 'gini = 0.549\nsamples = 18\nvalue = [3, 16, 8]\nnclass =
b'),
Text(2471.142857142857, 181.1999999999982, 'gini = 0.162\nsamples = 56\nvalue = [3, 84, 5]\nnclass =
a')]
```

```
b'),
Text(2869.7142857142853, 906.0, 'BEN <= 1.535\ngini = 0.638\nsamples = 1396\nvalue = [883, 434, 912]\n
class = c'),
Text(2710.285714285714, 543.5999999999999, 'CO <= 0.295\ngini = 0.571\nsamples = 825\nvalue = [267, 28
3, 776]\nnclass = c'),
Text(2630.5714285714284, 181.1999999999982, 'gini = 0.324\nsamples = 188\nvalue = [21, 248, 37]\nclas
s = b'),
Text(2790.0, 181.1999999999982, 'gini = 0.416\nsamples = 637\nvalue = [246, 35, 739]\nnclass = c'),
Text(3029.142857142857, 543.5999999999999, 'EBE <= 1.145\ngini = 0.484\nsamples = 571\nvalue = [616, 1
51, 136]\nnclass = a'),
Text(2949.428571428571, 181.1999999999982, 'gini = 0.211\nsamples = 17\nvalue = [3, 22, 0]\nnclass =
b'),
Text(3108.8571428571427, 181.1999999999982, 'gini = 0.467\nsamples = 554\nvalue = [613, 129, 136]\nnc
lass = a'),
Text(3826.2857142857138, 1268.4, 'BEN <= 1.915\ngini = 0.529\nsamples = 733\nvalue = [711, 91, 381]\nnc
lass = a'),
Text(3507.428571428571, 906.0, 'TOL <= 7.145\ngini = 0.538\nsamples = 221\nvalue = [94, 38, 207]\nnc
lass = c'),
Text(3347.999999999995, 543.5999999999999, 'NOx <= 105.85\ngini = 0.59\nsamples = 135\nvalue = [84, 2
3, 100]\nnclass = c'),
Text(3268.285714285714, 181.1999999999982, 'gini = 0.389\nsamples = 42\nvalue = [2, 14, 48]\nnclass =
c'),
Text(3427.7142857142853, 181.1999999999982, 'gini = 0.535\nsamples = 93\nvalue = [82, 9, 52]\nnclass =
a'),
Text(3666.8571428571427, 543.5999999999999, 'NOx <= 116.15\ngini = 0.324\nsamples = 86\nvalue = [10,
15, 107]\nnclass = c'),
Text(3587.142857142857, 181.1999999999982, 'gini = 0.531\nsamples = 6\nvalue = [1, 5, 2]\nnclass =
b'),
Text(3746.5714285714284, 181.1999999999982, 'gini = 0.271\nsamples = 80\nvalue = [9, 10, 105]\nnclass
= c'),
Text(4145.142857142857, 906.0, 'MXY <= 6.935\ngini = 0.419\nsamples = 512\nvalue = [617, 53, 174]\nnc
lass = a'),
Text(3985.7142857142853, 543.5999999999999, 'CO <= 0.46\ngini = 0.354\nsamples = 345\nvalue = [438, 3
3, 86]\nnclass = a'),
Text(3905.999999999995, 181.1999999999982, 'gini = 0.108\nsamples = 28\nvalue = [2, 33, 0]\nnclass =
b'),
Text(4065.428571428571, 181.1999999999982, 'gini = 0.275\nsamples = 317\nvalue = [436, 0, 86]\nnclass
= a'),
Text(4304.571428571428, 543.5999999999999, 'OXY <= 3.135\ngini = 0.512\nsamples = 167\nvalue = [179, 2
0, 88]\nnclass = a'),
Text(4224.857142857142, 181.1999999999982, 'gini = 0.567\nsamples = 45\nvalue = [16, 18, 49]\nnclass =
c'),
Text(4384.285714285714, 181.1999999999982, 'gini = 0.325\nsamples = 122\nvalue = [163, 2, 39]\nnclass
= a)']
```



Conclusion

Linear Regression=0.24039322600259538

Ridge Regression=0.23975318959885694

Lasso Regression=0.013157253655736056

ElasticNet Regression=0.08889326711786572

Logistic Regression=0.8101

Random Forest=0.8162857142857143

Random Forest is Suitable for this Dataset