

Vijay(P9) 3/08/2023

```
In [67]: 1 import numpy as np
          2 import pandas as pd
          3 import seaborn as sns
          4 import matplotlib.pyplot as plt
```

```
In [133]: 1 df=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\madrid_2009.csv")
          2 df
```

Out[133]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PM25 | PXY |
|--------|---------------------|------|------|------|------|------|-----------|------------|------|-----------|-----------|-------|------|
| 0 | 2009-10-01 01:00:00 | NaN | 0.27 | NaN | NaN | NaN | 39.889999 | 48.150002 | NaN | 50.680000 | 18.260000 | NaN | NaN |
| 1 | 2009-10-01 01:00:00 | NaN | 0.22 | NaN | NaN | NaN | 21.230000 | 24.260000 | NaN | 55.880001 | 10.580000 | NaN | NaN |
| 2 | 2009-10-01 01:00:00 | NaN | 0.18 | NaN | NaN | NaN | 31.230000 | 34.880001 | NaN | 49.060001 | 25.190001 | NaN | NaN |
| 3 | 2009-10-01 01:00:00 | 0.95 | 0.33 | 1.43 | 2.68 | 0.25 | 55.180000 | 81.360001 | 1.57 | 36.669998 | 26.530001 | 6.82 | 1.30 |
| 4 | 2009-10-01 01:00:00 | NaN | 0.41 | NaN | NaN | 0.12 | 61.349998 | 76.260002 | NaN | 38.090000 | 23.760000 | NaN | NaN |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 215683 | 2009-06-01 00:00:00 | 0.50 | 0.22 | 0.39 | 0.75 | 0.09 | 22.000000 | 24.510000 | 1.00 | 82.239998 | 10.830000 | 7.15 | 0.74 |
| 215684 | 2009-06-01 00:00:00 | NaN | 0.31 | NaN | NaN | NaN | 76.110001 | 101.099998 | NaN | 41.220001 | 9.920000 | NaN | NaN |
| 215685 | 2009-06-01 00:00:00 | 0.13 | NaN | 0.86 | NaN | 0.23 | 81.050003 | 99.849998 | NaN | 24.830000 | 12.460000 | 6.77 | NaN |
| 215686 | 2009-06-01 00:00:00 | 0.21 | NaN | 2.96 | NaN | 0.10 | 72.419998 | 82.959999 | NaN | NaN | 13.030000 | NaN | NaN |
| 215687 | 2009-06-01 00:00:00 | 0.37 | 0.32 | 0.99 | 1.36 | 0.14 | 54.290001 | 64.480003 | 1.06 | 56.919998 | 15.360000 | 11.61 | 0.83 |

215688 rows × 17 columns

◀ ▶

```
In [134]: 1 df=df.dropna()
```

```
In [135]: 1 df.columns
          2
```

```
Out[135]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
                 dtype='object')
```

In [136]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24717 entries, 3 to 215687
Data columns (total 17 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   date      24717 non-null   object  
 1   BEN        24717 non-null   float64 
 2   CO         24717 non-null   float64 
 3   EBE        24717 non-null   float64 
 4   MXY        24717 non-null   float64 
 5   NMHC       24717 non-null   float64 
 6   NO_2       24717 non-null   float64 
 7   NOx        24717 non-null   float64 
 8   OXY        24717 non-null   float64 
 9   O_3        24717 non-null   float64 
 10  PM10       24717 non-null   float64 
 11  PM25       24717 non-null   float64 
 12  PXY        24717 non-null   float64 
 13  SO_2       24717 non-null   float64 
 14  TCH        24717 non-null   float64 
 15  TOL        24717 non-null   float64 
 16  station    24717 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

In [137]: 1 data=df[['BEN', 'CO', 'station']]
2 data
3

Out[137]:

| | BEN | CO | station |
|--------|------|------|----------|
| 3 | 0.95 | 0.33 | 28079006 |
| 20 | 0.38 | 0.32 | 28079024 |
| 24 | 0.55 | 0.24 | 28079099 |
| 28 | 0.65 | 0.21 | 28079006 |
| 45 | 0.38 | 0.30 | 28079024 |
| ... | ... | ... | ... |
| 215659 | 0.54 | 0.27 | 28079024 |
| 215663 | 0.74 | 0.35 | 28079099 |
| 215667 | 0.78 | 0.29 | 28079006 |
| 215683 | 0.50 | 0.22 | 28079024 |
| 215687 | 0.37 | 0.32 | 28079099 |

24717 rows × 3 columns

```
In [138]: 1 df=df.head(15000)
2 df
```

Out[138]:

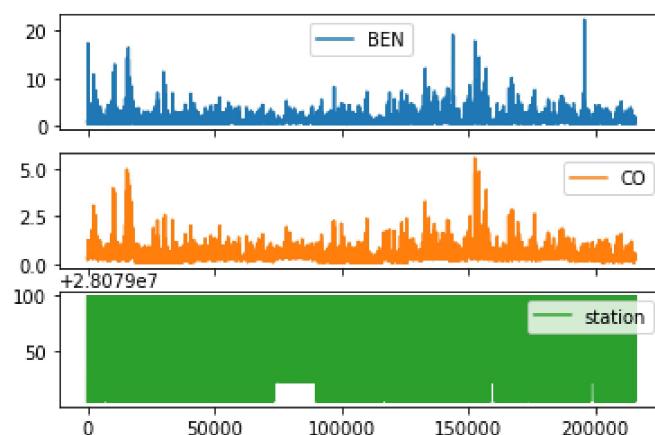
| | | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PM25 | PXY |
|--------|--|---------------------|------|------|------|------|------|------------|------------|------|-----------|-----------|-------|------|
| 3 | | 2009-10-01 01:00:00 | 0.95 | 0.33 | 1.43 | 2.68 | 0.25 | 55.180000 | 81.360001 | 1.57 | 36.669998 | 26.530001 | 6.82 | 1.30 |
| 20 | | 2009-10-01 01:00:00 | 0.38 | 0.32 | 0.32 | 0.89 | 0.01 | 17.969999 | 19.240000 | 1.00 | 65.870003 | 10.520000 | 7.01 | 0.84 |
| 24 | | 2009-10-01 01:00:00 | 0.55 | 0.24 | 0.65 | 1.79 | 0.18 | 36.619999 | 43.919998 | 1.28 | 48.070000 | 19.150000 | 9.33 | 1.07 |
| 28 | | 2009-10-01 02:00:00 | 0.65 | 0.21 | 1.20 | 2.04 | 0.18 | 37.169998 | 48.869999 | 1.21 | 26.950001 | 32.200001 | 6.94 | 1.04 |
| 45 | | 2009-10-01 02:00:00 | 0.38 | 0.30 | 0.50 | 1.15 | 0.00 | 17.889999 | 19.299999 | 1.00 | 60.009998 | 12.260000 | 8.46 | 0.88 |
| ... | | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 133346 | | 2009-02-12 23:00:00 | 2.17 | 0.69 | 1.91 | 3.75 | 0.24 | 93.320000 | 190.199997 | 2.14 | 7.150000 | 20.379999 | 11.89 | 1.77 |
| 133362 | | 2009-02-12 23:00:00 | 0.56 | 0.33 | 0.45 | 1.17 | 0.17 | 41.470001 | 44.639999 | 1.00 | 34.520000 | 8.750000 | 5.13 | 0.61 |
| 133366 | | 2009-02-12 23:00:00 | 0.72 | 0.55 | 0.96 | 2.46 | 0.19 | 80.019997 | 134.199997 | 1.57 | 13.970000 | 17.650000 | 9.23 | 1.19 |
| 133371 | | 2009-02-13 00:00:00 | 2.00 | 1.04 | 1.70 | 3.32 | 0.31 | 125.800003 | 302.399994 | 1.88 | 2.180000 | 32.790001 | 20.27 | 1.56 |
| 133387 | | 2009-02-13 00:00:00 | 0.64 | 0.55 | 0.45 | 1.16 | 0.18 | 67.540001 | 81.190002 | 1.00 | 9.880000 | 12.260000 | 7.16 | 0.60 |

15000 rows × 17 columns



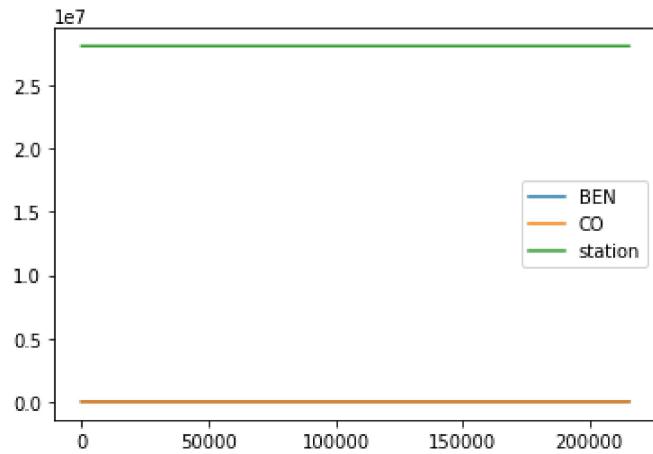
```
In [139]: 1 data.plot.line(subplots=True)
```

Out[139]: array([<AxesSubplot:>, <AxesSubplot:>, <AxesSubplot:>], dtype=object)



```
In [140]: 1 data.plot.line()  
2
```

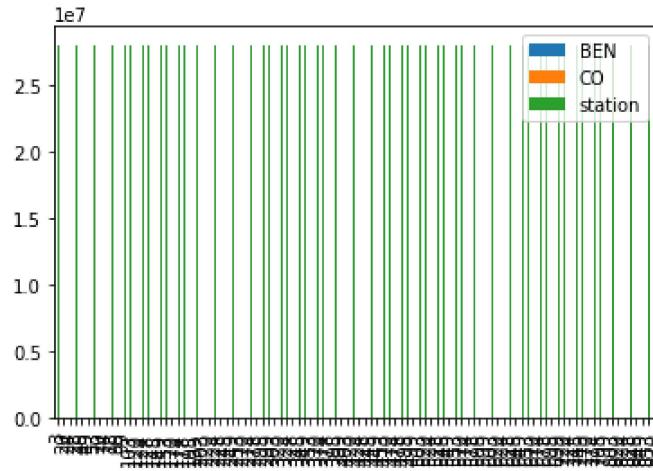
Out[140]: <AxesSubplot:>



```
In [143]: 1 b=data[0:100]  
2
```

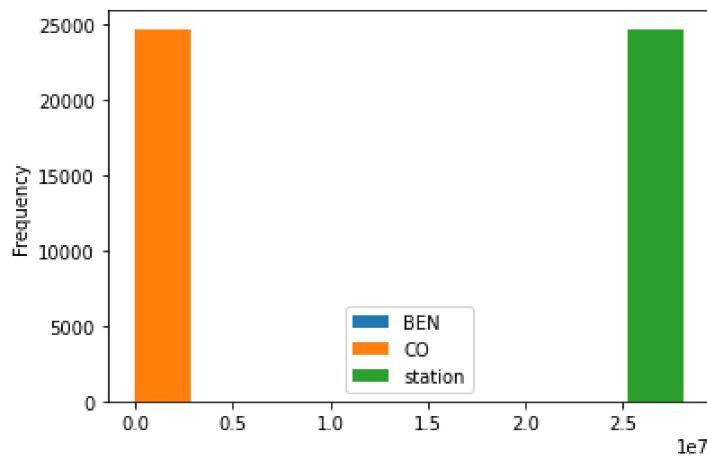
```
In [144]: 1 b.plot.bar()
```

Out[144]: <AxesSubplot:>



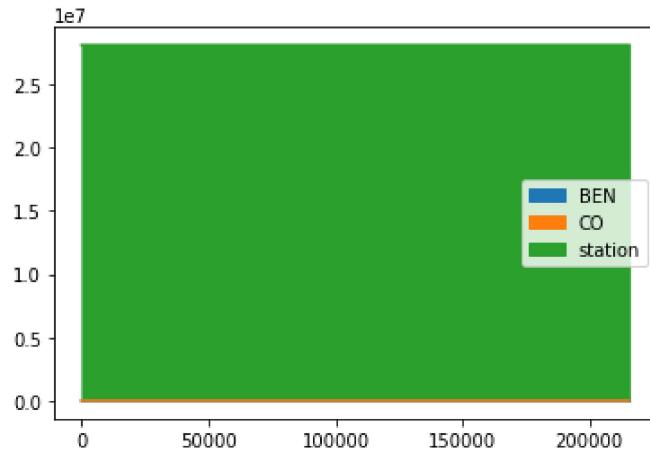
```
In [145]: 1 data.plot.hist()  
2
```

Out[145]: <AxesSubplot:ylabel='Frequency'>



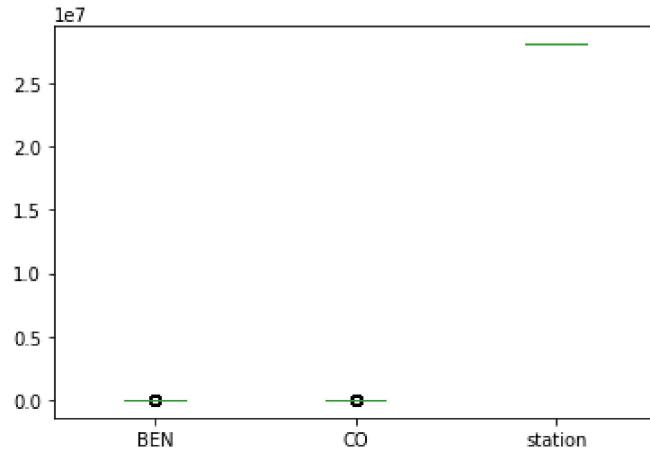
```
In [146]: 1 data.plot.area()
```

Out[146]: <AxesSubplot:>



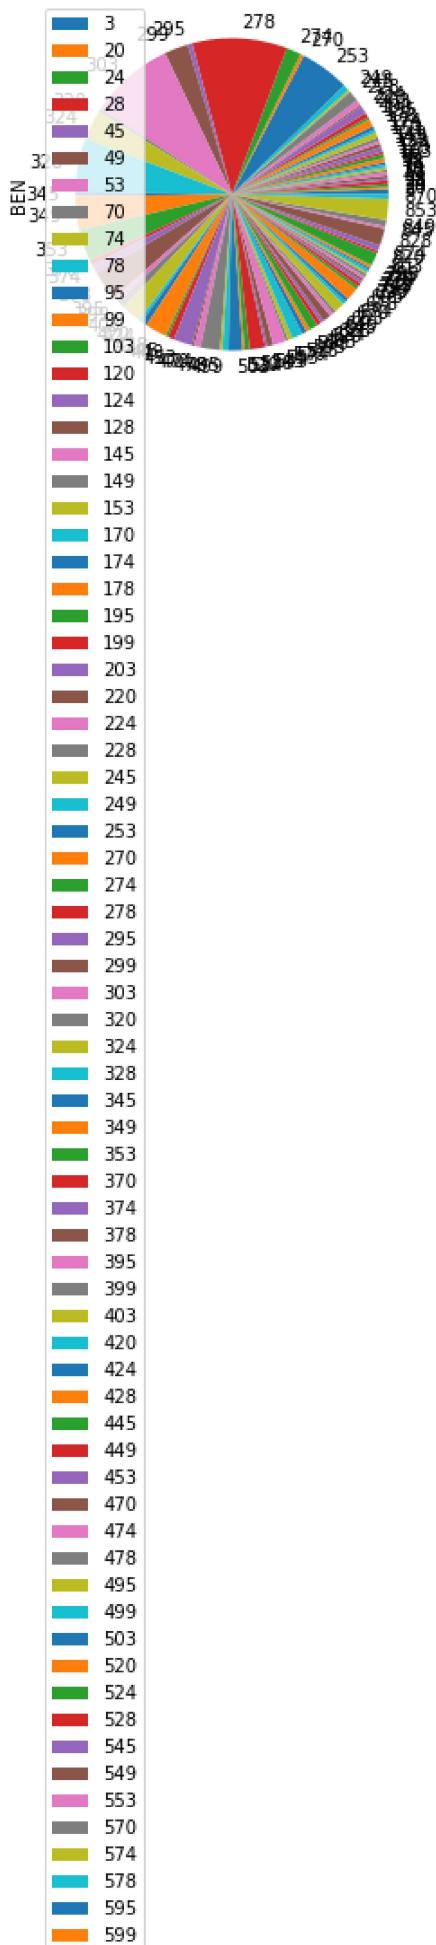
```
In [147]: 1 data.plot.box()
```

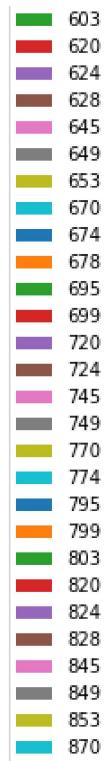
Out[147]: <AxesSubplot:>



```
In [148]: 1 b.plot.pie(y='BEN' )
```

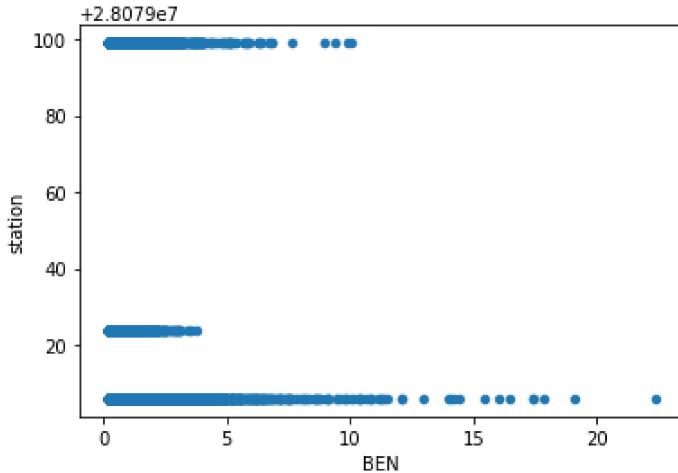
```
Out[148]: <AxesSubplot:ylabel='BEN'>
```



```
In [149]: 1 data.plot.scatter(x='BEN' ,y='station')  
2
```

```
Out[149]: <AxesSubplot:xlabel='BEN', ylabel='station'>
```



In [150]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 15000 entries, 3 to 133387
Data columns (total 17 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   date      15000 non-null   object  
 1   BEN        15000 non-null   float64 
 2   CO         15000 non-null   float64 
 3   EBE        15000 non-null   float64 
 4   MXY        15000 non-null   float64 
 5   NMHC       15000 non-null   float64 
 6   NO_2       15000 non-null   float64 
 7   NOx        15000 non-null   float64 
 8   OXY        15000 non-null   float64 
 9   O_3         15000 non-null   float64 
 10  PM10       15000 non-null   float64 
 11  PM25       15000 non-null   float64 
 12  PXY        15000 non-null   float64 
 13  SO_2       15000 non-null   float64 
 14  TCH         15000 non-null   float64 
 15  TOL         15000 non-null   float64 
 16  station    15000 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 2.1+ MB
```

In [151]: 1 df.describe()
2

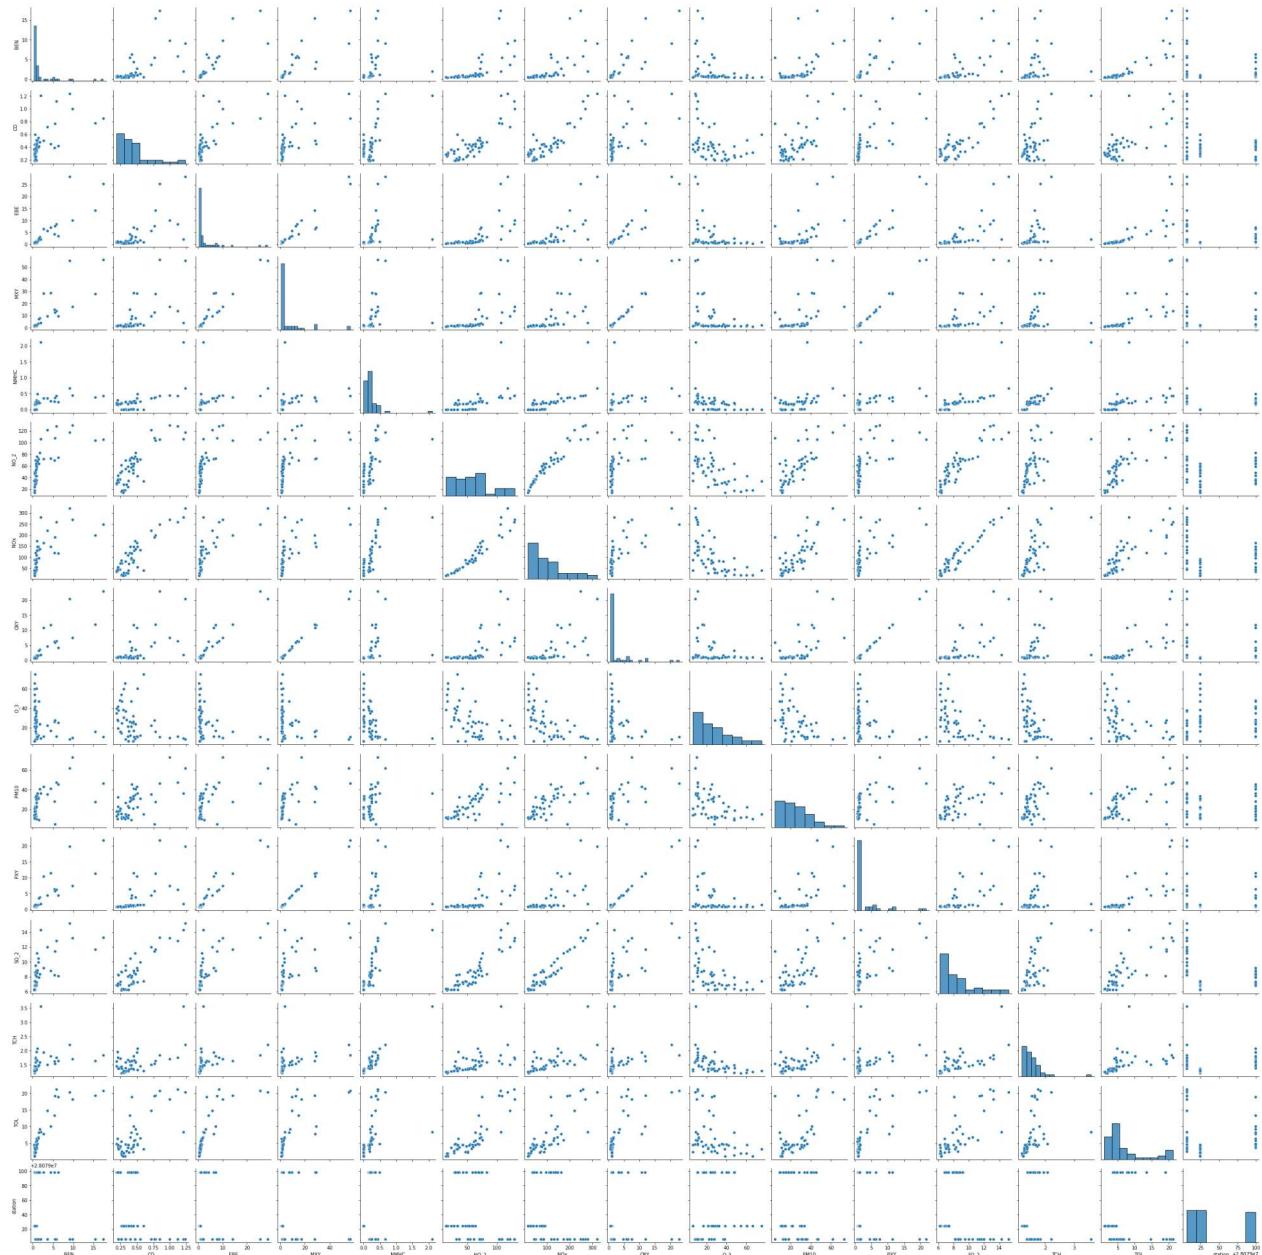
Out[151]:

| | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | |
|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|-------|
| count | 15000.000000 | 15000.000000 | 15000.000000 | 15000.000000 | 15000.000000 | 15000.000000 | 15000.000000 | 15000 |
| mean | 0.930545 | 0.424697 | 1.209391 | 2.115665 | 0.221291 | 54.027933 | 86.398438 | 1 |
| std | 0.937072 | 0.284678 | 0.942650 | 2.151884 | 0.152719 | 39.598022 | 87.947826 | 0 |
| min | 0.170000 | 0.060000 | 0.250000 | 0.240000 | 0.000000 | 0.860000 | 2.640000 | 0 |
| 25% | 0.440000 | 0.250000 | 0.720000 | 0.940000 | 0.140000 | 24.700001 | 30.510000 | 0 |
| 50% | 0.610000 | 0.350000 | 1.000000 | 1.350000 | 0.190000 | 45.184999 | 62.094999 | 1 |
| 75% | 1.060000 | 0.560000 | 1.370000 | 2.680000 | 0.250000 | 74.222502 | 115.300003 | 1 |
| max | 17.400000 | 4.980000 | 28.410000 | 56.500000 | 2.110000 | 477.399994 | 1438.000000 | 22 |

In [152]: 1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]

```
In [153]: 1 sns.pairplot(df1[0:50])
```

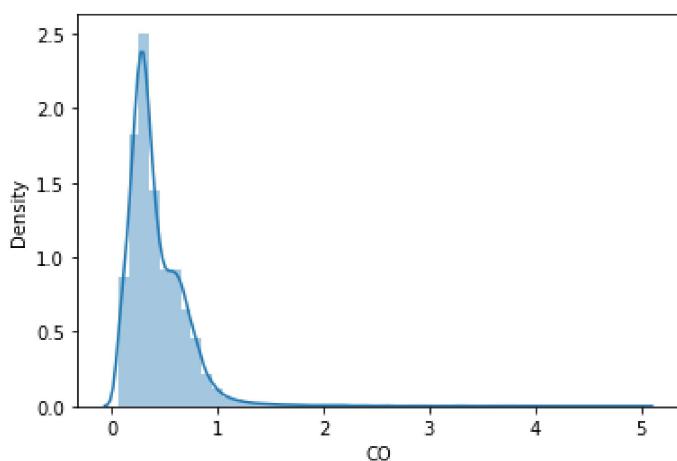
```
Out[153]: <seaborn.axisgrid.PairGrid at 0x27d329b75b0>
```



```
In [154]: 1 sns.distplot(df1['CO'])
2
```

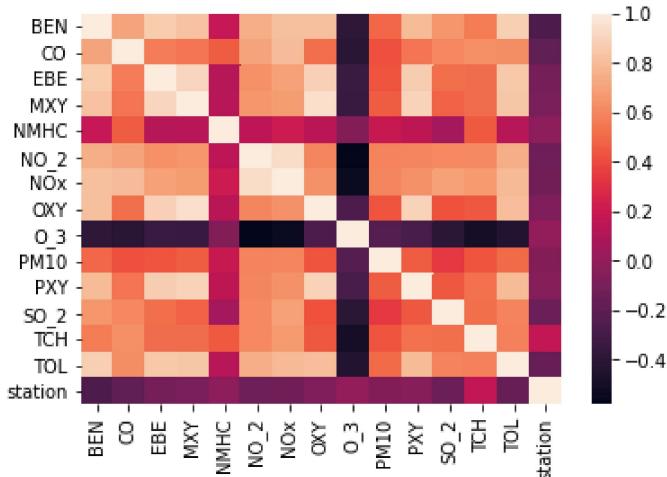
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
`warnings.warn(msg, FutureWarning)

```
Out[154]: <AxesSubplot:xlabel='CO', ylabel='Density'>
```



```
In [155]: 1 sns.heatmap(df1.corr())
```

```
Out[155]: <AxesSubplot:>
```



```
In [156]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2   'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
4
```

```
In [157]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
3
```

```
In [158]: 1 from sklearn.linear_model import LinearRegression
2 lr=LinearRegression()
3 lr.fit(x_train,y_train)
4
```

Out[158]: LinearRegression()

```
In [159]: 1 lr.intercept_
```

Out[159]: 28078835.04228549

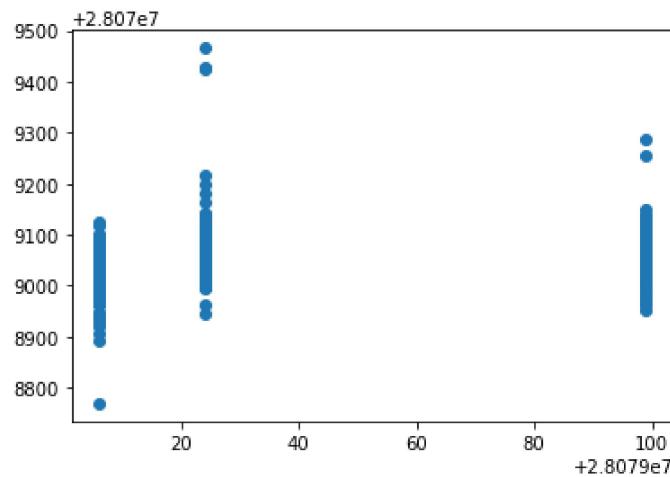
```
In [160]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
2 coeff
```

Out[160]:

| | Co-efficient |
|------|--------------|
| BEN | -32.629401 |
| CO | -26.427462 |
| EBC | 4.549725 |
| MXY | -4.091074 |
| NMHC | -46.142858 |
| NO_2 | -0.317408 |
| NOx | 0.201844 |
| OXY | 19.935779 |
| O_3 | 0.005785 |
| PM10 | -0.000663 |
| PXY | 2.293015 |
| SO_2 | -0.759300 |
| TCH | 176.605802 |
| TOL | -1.164715 |

```
In [161]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

Out[161]: <matplotlib.collections.PathCollection at 0x27d44a40c40>



```
In [162]: 1 lr.score(x_test,y_test)  
2
```

Out[162]: 0.2714591646166624

```
In [163]: 1 lr.score(x_train,y_train)  
2
```

Out[163]: 0.32874346114212316

```
In [164]: 1 from sklearn.linear_model import Ridge,Lasso  
2
```

```
In [165]: 1 rr=Ridge(alpha=10)  
2 rr.fit(x_train,y_train)  
3
```

Out[165]: Ridge(alpha=10)

```
In [166]: 1 rr.score(x_test,y_test)  
2
```

Out[166]: 0.28198376549370974

```
In [167]: 1 rr.score(x_train,y_train)  
2
```

Out[167]: 0.3265312248018226

```
In [168]: 1 la=Lasso(alpha=10)  
2 la.fit(x_train,y_train)
```

Out[168]: Lasso(alpha=10)

```
In [169]: 1 la.score(x_train,y_train)  
2
```

Out[169]: 0.03594355902967106

```
In [170]: 1 la.score(x_test,y_test)  
2
```

Out[170]: 0.02864461246144845

```
In [171]: 1 from sklearn.linear_model import ElasticNet  
2 en=ElasticNet()  
3 en.fit(x_train,y_train)  
4
```

Out[171]: ElasticNet()

```
In [172]: 1 en.coef_  
2
```

Out[172]: array([-6.21414103, -0.88172107, 0. , 1.58857296, -0. ,
-0.32394856, 0.13688405, 1.00189763, -0.16854009, 0.12559179,
2.40741098, -0.96157604, 1.75509927, -1.66205067])

```
In [173]: 1 en.intercept_
2
```

Out[173]: 28079068.769639283

```
In [174]: 1 prediction=en.predict(x_test)
2
```

```
In [175]: 1 en.score(x_test,y_test)
```

Out[175]: 0.09100787558661116

```
In [176]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
5
```

36.393475561408536

1502.6854920865546

38.76448751223928

```
In [177]: 1 from sklearn.linear_model import LogisticRegression
2
```

```
In [178]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df['station']
4
```

```
In [179]: 1 feature_matrix.shape
2
```

Out[179]: (15000, 14)

```
In [180]: 1 target_vector.shape
2
```

Out[180]: (15000,)

```
In [181]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [182]: 1 fs=StandardScaler().fit_transform(feature_matrix)
2
```

```
In [183]: 1 logr=LogisticRegression(max_iter=10000)
2 logr.fit(fs,target_vector)
```

Out[183]: LogisticRegression(max_iter=10000)

```
In [184]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
2
```

```
In [185]: 1 prediction=logr.predict(observation)
           2 print(prediction)
           3
```

```
[28079099]
```

```
In [186]: 1 logr.classes_
           2
```

```
Out[186]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [187]: 1 logr.score(fs,target_vector)
           2
```

```
Out[187]: 0.9266
```

```
In [188]: 1 logr.predict_proba(observation)[0][0]
           2
```

```
Out[188]: 1.741833318862e-13
```

```
In [189]: 1 logr.predict_proba(observation)
           2
```

```
Out[189]: array([[1.74183332e-13, 2.94598015e-42, 1.00000000e+00]])
```

```
In [190]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [191]: 1 rfc=RandomForestClassifier()
           2 rfc.fit(x_train,y_train)
```

```
Out[191]: RandomForestClassifier()
```

```
In [192]: 1 parameters={'max_depth':[1,2,3,4,5],
           2   'min_samples_leaf':[5,10,15,20,25],
           3   'n_estimators':[10,20,30,40,50]
           4 }
```

```
In [193]: 1 from sklearn.model_selection import GridSearchCV
           2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
           3 grid_search.fit(x_train,y_train)
           4
```

```
Out[193]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                       param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                       scoring='accuracy')
```

```
In [194]: 1 grid_search.best_score_
```

```
Out[194]: 0.9172380952380952
```

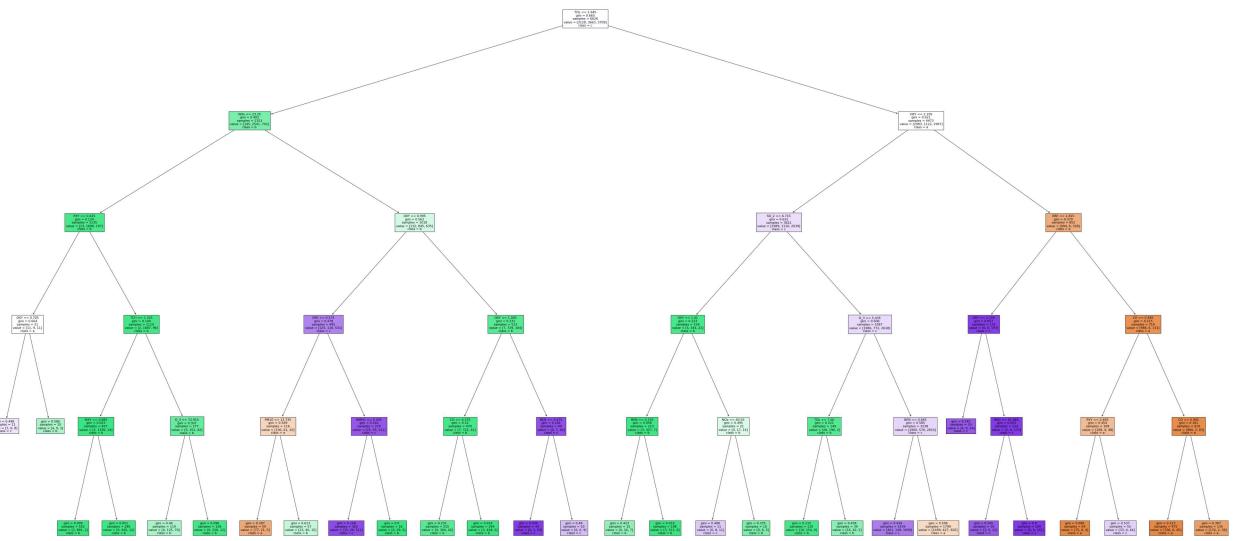
```
In [195]: 1 rfc_best=grid_search.best_estimator_
```

In [196]:

```
1 from sklearn.tree import plot_tree
2 plt.figure(figsize=(80,40))
3 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'])
```

```
Out[196]: [Text(2150.8363636363633, 1993.2, 'TOL <= 1.645\nngini = 0.665\nsamples = 6626\nvalue = [31
28, 3663, 3709]\nclass = c'),
Text(953.6727272727272, 1630.8000000000002, 'NOx <= 23.25\nngini = 0.402\nsamples = 2153\nvalue = [145, 2541, 742]\nclass = b'),
Text(365.2363636363636, 1268.4, 'PXY <= 0.425\nngini = 0.124\nsamples = 1135\nvalue = [13,
1696, 107]\nclass = b'),
Text(162.3272727272727, 906.0, 'OXY <= 0.705\nngini = 0.664\nsamples = 21\nvalue = [11, 9,
11]\nclass = a'),
Text(81.16363636363636, 543.5999999999999, 'gini = 0.498\nsamples = 11\nvalue = [7, 0, 8]
\nclass = c'),
Text(243.49090909090907, 543.5999999999999, 'gini = 0.586\nsamples = 10\nvalue = [4, 9,
3]\nclass = b'),
Text(568.1454545454545, 906.0, 'TCH <= 1.325\nngini = 0.104\nsamples = 1114\nvalue = [2, 1
687, 96]\nclass = b'),
Text(405.818181818176, 543.5999999999999, 'MXY <= 0.885\nngini = 0.023\nsamples = 837\nvalue = [2, 1336, 14]\nclass = b'),
Text(324.6545454545454, 181.1999999999982, 'gini = 0.009\nsamples = 552\nvalue = [2, 89
1, 2]\nclass = b'),
Text(486.981818181814, 181.1999999999982, 'gini = 0.051\nsamples = 285\nvalue = [0, 44
5, 12]\nclass = b'),
Text(730.4727272727272, 543.5999999999999, 'O_3 <= 72.915\nngini = 0.307\nsamples = 277\nvalue = [0, 351, 82]\nclass = b'),
Text(649.3090909090909, 181.1999999999982, 'gini = 0.46\nsamples = 119\nvalue = [0, 125,
70]\nclass = b'),
Text(811.6363636363635, 181.1999999999982, 'gini = 0.096\nsamples = 158\nvalue = [0, 22
6, 12]\nclass = b'),
Text(1542.1090909090908, 1268.4, 'OXY <= 0.995\nngini = 0.563\nsamples = 1018\nvalue = [13
2, 845, 635]\nclass = b'),
Text(1217.45454545453, 906.0, 'EBE <= 0.575\nngini = 0.478\nsamples = 495\nvalue = [125,
116, 531]\nclass = c'),
Text(1055.1272727272726, 543.5999999999999, 'PM10 <= 13.735\nngini = 0.569\nsamples = 116
\nvalue = [100, 61, 20]\nclass = a'),
Text(973.9636363636363, 181.1999999999982, 'gini = 0.397\nsamples = 59\nvalue = [77, 21,
5]\nclass = a'),
Text(1136.290909090909, 181.1999999999982, 'gini = 0.613\nsamples = 57\nvalue = [23, 40,
15]\nclass = b'),
Text(1379.78181818182, 543.5999999999999, 'NMHC <= 0.245\nngini = 0.242\nsamples = 379\nvalue = [25, 55, 511]\nclass = c'),
Text(1298.61818181817, 181.1999999999982, 'gini = 0.169\nsamples = 363\nvalue = [25, 2
6, 511]\nclass = c'),
Text(1460.9454545454544, 181.1999999999982, 'gini = 0.0\nsamples = 16\nvalue = [0, 29,
0]\nclass = b'),
Text(1866.7636363636361, 906.0, 'OXY <= 1.005\nngini = 0.231\nsamples = 523\nvalue = [7, 7
29, 104]\nclass = b'),
Text(1704.4363636363635, 543.5999999999999, 'CO <= 0.375\nngini = 0.12\nsamples = 479\nvalue = [7, 722, 42]\nclass = b'),
Text(1623.272727272727, 181.1999999999982, 'gini = 0.231\nsamples = 215\nvalue = [4, 30
4, 42]\nclass = b'),
Text(1785.6, 181.1999999999982, 'gini = 0.014\nsamples = 264\nvalue = [3, 418, 0]\nclass
= b'),
Text(2029.090909090909, 543.5999999999999, 'BEN <= 0.475\nngini = 0.182\nsamples = 44\nvalue = [0, 7, 62]\nclass = c'),
Text(1947.9272727272726, 181.1999999999982, 'gini = 0.036\nsamples = 34\nvalue = [0, 1,
53]\nclass = c'),
Text(2110.254545454545, 181.1999999999982, 'gini = 0.48\nsamples = 10\nvalue = [0, 6, 9]
\nclass = c'),
Text(3347.999999999995, 1630.8000000000002, 'OXY <= 2.105\nngini = 0.621\nsamples = 4473
\nvalue = [2983, 1122, 2967]\nclass = a'),
Text(2840.7272727272725, 1268.4, 'SO_2 <= 6.715\nngini = 0.631\nsamples = 3621\nvalue = [1
989, 1116, 2639]\nclass = c'),
Text(2516.072727272727, 906.0, 'OXY <= 1.01\nngini = 0.123\nsamples = 234\nvalue = [3, 34
4, 21]\nclass = b'),
Text(2353.7454545454543, 543.5999999999999, 'BEN <= 0.335\nngini = 0.058\nsamples = 213\nnv')]
```

```
alue = [3, 327, 7]\nclass = b'),  
Text(2272.581818181818, 181.19999999999982, 'gini = 0.423\nsamples = 15\nvalue = [0, 16,  
7]\nclass = b'),  
Text(2434.9090909090905, 181.19999999999982, 'gini = 0.019\nsamples = 198\nvalue = [3, 31  
1, 0]\nclass = b'),  
Text(2678.3999999999996, 543.5999999999999, 'NOx <= 40.55\ngini = 0.495\nsamples = 21\nva  
lue = [0, 17, 14]\nclass = b'),  
Text(2597.2363636363634, 181.19999999999982, 'gini = 0.488\nsamples = 11\nvalue = [0, 8,  
11]\nclass = c'),  
Text(2759.5636363636363, 181.19999999999982, 'gini = 0.375\nsamples = 10\nvalue = [0, 9,  
3]\nclass = b'),  
Text(3165.381818181818, 906.0, 'O_3 <= 5.435\ngini = 0.606\nsamples = 3387\nvalue = [198  
6, 772, 2618]\nclass = c'),  
Text(3003.0545454545454, 543.5999999999999, 'TOL <= 7.06\ngini = 0.221\nsamples = 149\nva  
lue = [26, 196, 2]\nclass = b'),  
Text(2921.8909090909087, 181.19999999999982, 'gini = 0.115\nsamples = 110\nvalue = [10, 1  
54, 0]\nclass = b'),  
Text(3084.2181818181816, 181.19999999999982, 'gini = 0.438\nsamples = 39\nvalue = [16, 4  
2, 2]\nclass = b'),  
Text(3327.7090909090907, 543.5999999999999, 'BEN <= 0.665\ngini = 0.585\nsamples = 3238\n  
value = [1960, 576, 2616]\nclass = c'),  
Text(3246.545454545454, 181.19999999999982, 'gini = 0.416\nsamples = 1439\nvalue = [461,  
149, 1690]\nclass = c'),  
Text(3408.872727272727, 181.19999999999982, 'gini = 0.596\nsamples = 1799\nvalue = [1499,  
427, 926]\nclass = a'),  
Text(3855.272727272727, 1268.4, 'EBE <= 1.815\ngini = 0.379\nsamples = 852\nvalue = [994,  
6, 328]\nclass = a'),  
Text(3571.2, 906.0, 'OXY <= 2.155\ngini = 0.057\nsamples = 133\nvalue = [6, 0, 197]\nclas  
s = c'),  
Text(3490.036363636363, 543.5999999999999, 'gini = 0.245\nsamples = 19\nvalue = [4, 0, 2  
4]\nclass = c'),  
Text(3652.363636363636, 543.5999999999999, 'NOx <= 61.965\ngini = 0.023\nsamples = 114\nva  
lue = [2, 0, 173]\nclass = c'),  
Text(3571.2, 181.19999999999982, 'gini = 0.245\nsamples = 10\nvalue = [2, 0, 12]\nclass =  
c'),  
Text(3733.5272727272722, 181.19999999999982, 'gini = 0.0\nsamples = 104\nvalue = [0, 0, 1  
61]\nclass = c'),  
Text(4139.345454545454, 906.0, 'CO <= 0.485\ngini = 0.215\nsamples = 719\nvalue = [988,  
6, 131]\nclass = a'),  
Text(3977.0181818181813, 543.5999999999999, 'PXY <= 2.405\ngini = 0.454\nsamples = 109\nva  
lue = [108, 4, 48]\nclass = a'),  
Text(3895.854545454545, 181.19999999999982, 'gini = 0.096\nsamples = 54\nvalue = [75, 0,  
4]\nclass = a'),  
Text(4058.181818181818, 181.19999999999982, 'gini = 0.537\nsamples = 55\nvalue = [33, 4,  
44]\nclass = c'),  
Text(4301.672727272727, 543.5999999999999, 'CO <= 0.905\ngini = 0.161\nsamples = 610\nva  
lue = [880, 2, 83]\nclass = a'),  
Text(4220.50909090909, 181.19999999999982, 'gini = 0.113\nsamples = 475\nvalue = [706, 0,  
45]\nclass = a'),  
Text(4382.836363636363, 181.19999999999982, 'gini = 0.307\nsamples = 135\nvalue = [174,  
2, 38]\nclass = a')
```



Conclusion

Linear Regression=0.32874346114212316

Ridge Regression=0.3265312248018226

Lasso Regression=0.03594355902967106

ElasticNet Regression=0.09100787558661116

Logistic Regression=0.9266

Random Forest=0.9172380952380952

Logistic Regression is Suitable for this Dataset